

Machine Learning

Lecture 2: k -Nearest Neighbors

Prof. Dr. Stephan Günnemann

Data Mining and Analytics
Technical University of Munich

15.10.2019

Main reading

- "Machine Learning: A Probabilistic Perspective" by Murphy [ch. 1.4.1 - 1.4.3]

Extra reading

- "Bayesian Reasoning and Machine Learning" by Barber [ch. 14]

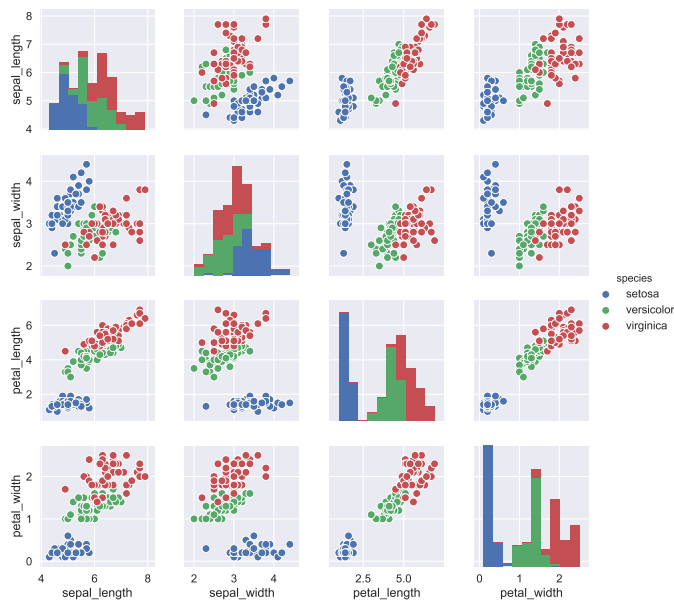
Slides adapted from previous versions by W. Koepp & D. Korhammer

k-Nearest Neighbors

1

Data Mining
and Analytics TUM

Iris dataset



Iris dataset: 2 features




How do we intuitively label new samples by hand?
Look at the *surrounding* points. Do as your *neighbor* does.

Given a training dataset $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$
where $\mathbf{x}_i \in \mathbb{R}^D$ are features and $y_i \in \{1, \dots, C\}$ are class labels

To classify new observations:

- define a distance measure (e.g. Euclidean distance)
- compute the nearest neighbor for all new data points
- and label them with the label of their nearest neighbor

This works for both *classification* and *regression*.



This corresponds to a Voronoi tessellation.
And results in poor generalization...

k-Nearest Neighbor classification

More *robust* against errors in the training set:

Look at multiple nearest neighbors and pick the **majority** label.

Let $\mathcal{N}_k(\mathbf{x})$ be the k nearest neighbors of a vector \mathbf{x} , then in classification tasks:

$$p(y = c \mid \mathbf{x}, k) = \frac{1}{k} \sum_{i \in \mathcal{N}_k(\mathbf{x})} \mathbb{I}(y_i = c),$$
$$\hat{y} = \arg \max_c p(y = c \mid \mathbf{x}, k)$$

with the *indicator variable* $\mathbb{I}(e)$ is defined as:

$$\mathbb{I}(e) = \begin{cases} 1 & \text{if } e \text{ is true} \\ 0 & \text{if } e \text{ is false.} \end{cases}$$

i.e., the vector will be labeled by the mode of its neighbors' labels.

k-Nearest Neighbor classification: weighted

Look at multiple nearest neighbors and pick the **weighted majority** label.
The weight is **inversely proportional** to the distance.

Let $\mathcal{N}_k(\mathbf{x})$ be the k nearest neighbors of a vector \mathbf{x} , then in classification tasks:

$$p(y = c \mid \mathbf{x}, k) = \frac{1}{Z} \sum_{i \in \mathcal{N}_k(\mathbf{x})} \frac{1}{d(\mathbf{x}, \mathbf{x}_i)} \mathbb{I}(y_i = c),$$
$$\hat{y} = \arg \max_c p(y = c \mid \mathbf{x}, k)$$

with $Z = \sum_{i \in \mathcal{N}_k(\mathbf{x})} \frac{1}{d(\mathbf{x}, \mathbf{x}_i)}$ the normalization constant and $d(\mathbf{x}, \mathbf{x}_i)$ being a distance measure between \mathbf{x} and \mathbf{x}_i .

Regression is similar:


Let $\mathcal{N}_k(\mathbf{x})$ be the k nearest neighbors of a vector \mathbf{x} , then for regression:

$$\hat{y} = \frac{1}{Z} \sum_{i \in \mathcal{N}_k(\mathbf{x})} \frac{1}{d(\mathbf{x}, \mathbf{x}_i)} y_i,$$


with $Z = \sum_{i \in \mathcal{N}_k(\mathbf{x})} \frac{1}{d(\mathbf{x}, \mathbf{x}_i)}$ the normalization constant and $d(\mathbf{x}, \mathbf{x}_i)$ being a distance measure between \mathbf{x} and \mathbf{x}_i ,

i.e., the vector will be labeled by a **weighted mean** of its neighbors' values.

Note: y_i is a real number here (rather than categorical label).




So, how many neighbors are best?



Compare the decision boundaries of 1-NN and 3-NN

Goal is **generalization**: pick k (called a *hyper-parameter*) that performs best¹ on unseen (future) data.

Unfortunately, no access to future data, so split the dataset \mathcal{D} :




Hyper-parameter tuning procedure

- Learn the model using the training set
- Evaluate performance with different k on the *validation set* picking the best k
- Report final performance on the test set.²

¹In terms of some predefined metric, i.e. accuracy
²Good data science practices: See slides on Decision Trees

Using validation set to choose k



We choose $k = 7$.

Measuring classification performance

How can we assess the performance of a (binary) classification algorithm?

⇒ Confusion table


Predicted	True condition	
	$y = 1$	$y = 0$
$y = 1$	TP	FP
$y = 0$	FN	TN

- TP

TN

= true positive
= true negative
- FP

FN
- = false positive
= false negative
- } correct predictions
- } wrong predictions



Measuring classification performance

Accuracy:	$acc = \frac{TP + TN}{TP + TN + FP + FN}$
Precision:	$prec = \frac{TP}{TP + FP}$
Sensitivity/Recall:	$rec = \frac{TP}{TP + FN}$
Specificity:	$tnr = \frac{TN}{FP + TN}$
False Negative Rate:	$fnr = \frac{FN}{TP + FN}$
False Positive Rate:	$fpr = \frac{FP}{FP + TN}$
F1 Score:	$f1 = \frac{2 \cdot prec \cdot rec}{prec + rec}$

⇒ Trade-off between precision and recall: increasing one (most often) leads to decreasing the other


General note: Be careful when you have imbalanced classes!

Distance measures

- K-NN can be used with various distance measures → highly flexible
- Euclidean distance (L_2 norm): $\sqrt{\sum_i (u_i - v_i)^2}$

- L_1 norm: $\sum_i |u_i - v_i|$
- L_∞ norm: $\max_i |u_i - v_i|$
- Angle:


$$\cos \alpha = \frac{\mathbf{u}^T \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$



- Mahalanobis distance (Σ is positive (semi) definite and symmetric):

$$\sqrt{(\mathbf{u} - \mathbf{v})^T \Sigma^{-1} (\mathbf{u} - \mathbf{v})}$$

- Hamming distance, Edit distance, ...



The same old example but one of our features is in the order of meters, the other in the order of centimeters. ($k = 1$)

- Data *standardization*
Scale each feature to zero mean and unit variance.

$$x_{i,\text{std}} = \frac{x_i - \mu_i}{\sigma_i}$$

(This is a standard procedure in machine learning. Many models are sensitive to differences in scale.)

- Use the Mahalanobis distance.

$$\text{mahalanobis}(x_1, x_2) = \sqrt{(x_1 - x_2)^T \Sigma^{-1} (x_1 - x_2)}$$

$$\Sigma = \begin{bmatrix} \sigma_1^2 & 0 & 0 \\ 0 & \dots & 0 \\ 0 & 0 & \sigma_n^2 \end{bmatrix} \text{ is equal to Euclidean distance on normalized data}$$


The curse of dimensionality

Given a discrete one-dimensional input space $x \in \{1, 2, \dots, 10\}$

For $N = 20$ uniformly distributed samples the data covers 100% of the input space.


Add a second dimension (now $x \in \{1, \dots, 10\}^2$) and your data only covers 18% of the input space.

Once you add a third dimension you only cover 2%.



The curse of dimensionality


- The nearest neighbor will now be pretty far away..
- N has to grow exponentially with the number of features. Consider this when using k -NN on high-dimensional data.



Expensive: memory and naive inference are both $O(N)$:

we need to store the entire training data and compare with all training instances to find the nearest neighbor

Solution: use tree-based search structures (e.g. k-d tree) for efficient (approximate) NN³



³At the expense of an additional computation performed only once

- k -NN Algorithm
- Train-validation-test split
- Measuring classification performance
- Distance metrics
- Curse of dimensionality